

CHAPTER ONE



Two Clinicians, One Client

How identical assessment data produces two defensible plans — and why the difference isn't ABA knowledge.

THE PACKET

Mateo is three years and two months old. He was diagnosed sixteen weeks ago, and this is his first authorization. What follows is the intake packet, condensed to what both clinicians' in this chapter actually had. Read it the way you would read your own referral: not to memorize it, but to notice what you start doing with it.

Referral. Developmental pediatrician, following autism diagnosis at 2 years, 10 months. Referral line: "Comprehensive ABA evaluation and treatment; parents report communication delay and behavioral concerns." Recommended intensity from the diagnostic report: 25 hours weekly.

Skills assessment (direct testing and observation, two sessions). No vocal mands. Requests are made by pulling an adult by the wrist toward the desired item, roughly a dozen times per hour when motivated; occasional reaching and whining directed at items, not at people. Vocalizations are frequent but noncontingent — strings of vowel-heavy babble, two or three consonant sounds heard across both sessions. Echoic testing: no reliable point-to-point correspondence; he looked up at the examiner's face twice during vocal models. Motor imitation: three of ten gross-motor targets, inconsistent across trials. Responds to his name on roughly one attempt in five, and mostly when the caller is between him and something he wants. Matching: identical objects emerging, two of five arrays. Listener responding: touches two body parts, inconsistently; no object selection from an array. Play: dumps bins, mouths pieces, and lines up vehicles by size; the line, once built, is defended. He tolerated the table for stretches of two to four minutes before sliding off his chair; when materials were re-presented after the second slide, he pushed them off the table.

Preference and observation notes. Reliable reinforcers: the iPad (strongest by a wide margin), rough-and-tumble play, chase, water play at the sink, one specific brand of crackers. He brought his mother the iPad eleven times across one home observation. When the iPad was removed, tantrums followed on four of five occasions — crying, dropping to the floor, once throwing the case. Tantrums are otherwise brief (under three minutes), occur six to eight times daily, and cluster around demands and around the ends of preferred activities.

Caregiver interview. Mother, asked for her single most important hope for this year: "I want him to say Mama. And to look at me when I call him." Fathers priorities: the car seat (daily physical resistance, an escalating safety concern on the highway) and haircuts, which the family has abandoned entirely — Mateos hair has not

been cut in seven months. Sleep onset takes forty-five to ninety minutes with a parent lying beside him; he wakes for the day between 5:00 and 5:30. He eats approximately twelve foods, weight tracking along his curve, pediatrician aware and unconcerned for now. He attends a daycare program three mornings weekly, where staff describe him as "in his own world, but no trouble." No seizure history, no hearing test since the newborn screen. Both parents work; grandmother covers two afternoons.

That is the packet. Two board-certified behavior analysts received it in the same week — a quirk of a funding transition, and the reason this chapter exists. Each wrote a first-authorization treatment plan. Neither saw the others work. Both plans were reviewed by the same senior clinician, who judged both defensible, and who kept both, because together they teach something neither could teach alone.

You will read both plans shortly. Not yet.

YOUR TURN COMES FIRST

Here is the first rule of this book, and it is not negotiable: you commit before you see. The entire method of this text depends on it. Once you have read another clinician's plan, your own reasoning becomes unrecoverable — hindsight will quietly rewrite what you would have done into what you now know, and the gap between the two, which is the only thing this book can actually teach into, disappears. Every anchor case in this book will stop you at the fork and ask you to write. This is the first fork.

Before reading further, write your first-authorization plan for Mateo. Not a sketch — the real thing, at the level of specificity you would submit:

Your goals, by domain, in the order you would begin teaching them. Use Goal Bank IDs if you know them; plain language if you don't.

Your recommended weekly hours, and one sentence justifying the number.

Your answer to the mother. She asked for "Mama" and for his name. What does your plan say to her, and when will you say it?

Your postponement list: at least three goals you considered and deliberately did not write, each with one sentence of reasoning.

Three predictions, each with a percentage: the probability that Mateo emits his first functional vocal mand within this authorization period; the probability that your plan survives to reauthorization without a major revision; the probability that the car seat problem is resolved within six months.

Date the page. You will need it twice more in this book — once in Chapter 20, when your predictions come due, and once in Chapter 21, when a senior clinician who disagrees with you reviews it on the page. Write it somewhere you will not lose.

If you are tempted to skip this — you are tired, you are reading on a train, you will "do it mentally" — know that the mental version is worthless in a specific way: unwritten plans are unfalsifiable. You cannot be surprised by your own reasoning unless it exists somewhere your memory cannot edit. The clinicians' who get the most from

this book are the ones who let it embarrass them a little.

PLAN ONE — RAMOS

Ramos' has nine years of experience, most of it in early intervention. Her plan runs eleven pages; here is its goal architecture, verbatim in structure:

Hours requested: 25. Justification: severity across all domains, diagnostic recommendation, family capacity to support home programming.

Notice what Ramos' saw. She caught the sleep data and took it seriously — forty-five to ninety minutes of sleep onset with a parent trapped in the room is not a footnote; it is a tax collected nightly against every other goal in the plan, paid by the two people expected to implement it. She caught the feeding numbers and made a disciplined call: twelve foods with a stable growth curve is a watch-item, not a program, so she wrote surveillance instead of intervention — a log, not a hierarchy. That is restraint, and it is the correct kind. She addressed the iPad directly, and she anchored the whole plan in manding and imitation, which is where it should be anchored.

Her plan reads as thorough. Ten goals; every deficit the assessment surfaced now has a program with its name on it. A reviewer skimming it sees comprehensiveness. A family reading it sees that everything they mentioned was heard. It is the plan most of us were trained to write, and there are years in every career when it is the plan we did write.

Hold your objections. First read the other one.

PLAN TWO — WHITFIELD

Whitfield' has fourteen years of experience and a habit that shows up in his documents: every plan he writes has two lists, and the second one is what he didn't write.

Hours requested: 25 — the identical number, defended differently, and we will come back to how.

Whitfield's postponement list, from the plan itself:

Echoics (C1) — postponed 90 days. Echoic control is currently near zero, and drill-based echoic training under weak stimulus control in month one buys little language at the price of making the table aversive during the exact window when the table must be paired with reinforcement. Mand training under strong motivation is the better recruiter of vocal behavior for this profile: the shaping happens inside requests he already wants to make. Echoic status re-probed at day 90; C1 enters the plan the moment point-to-point correspondence emerges or the mand modality stalls.

Listener responding (B7) — postponed. Gated behind B11. An array-selection program for a child who orients to his name one time in five is a program about the technicians persistence, not the child's learning. Name first, then listener targets ride on it.

Sleep (R-domain) — deferred to medical review, then reconsidered. Sleep-onset duration plus 5 a.m. waking plus no audiological follow-up since birth earns a pediatrician conversation before a behavioral program. Referral letter attached. If medical review is unremarkable, R1 is the first candidate for the 90-day plan revision — and the family has been told exactly that.

Feeding (G8) — watch only. Same call as it happens Ramos' made, same reasoning.

Haircuts and car seat (O-domain / K13) — acknowledged, sequenced second. The car seat is a genuine safety item and the first behavior-reduction candidate of authorization two, after the transition and tolerance machinery built in goals 3 and 6 gives it something to stand on. The family has this in writing, with the sequencing reason.

Six goals. One page of programs, one page of reasons for the programs that aren't there.

THE COMPARISON — WHERE PLANS COME FROM

The claim this chapter makes, and the book will spend twenty more defending: these two plans do not differ because one clinician knows more ABA. Test that yourself — there is no fact about behavior analysis that Ramos' is missing. Both clinicians' saw the same wrist-pull and knew it was already a mand, needing transfer to a topography the world can read, not creation from nothing. Both saw that the iPad is simultaneously the strongest reinforcer in the house and the strongest trigger, which is not a contradiction but a single fact wearing two costumes. Both know the Early Learner presentation cold.

The plans differ because the two clinicians' gave different answers to a small set of questions that neither of them wrote down — questions that were answered, in Ramos's case, implicitly, and in Whitfield's case, on paper. Reconstructing those answers is the work of this section, and learning to reconstruct them — in colleagues' plans, in supervisees' plans, and above all in your own — is the first skill of clinical reasoning.

Question one: what does this authorization have to accomplish — everything it can, or the most it can? Ramos' answered: cover the assessed deficits; ten goals, each defensible on its own. Whitfield' answered: concentrate force. And here the arithmetic stops being an opinion. Twenty-five hours a week, spread across ten programs, is roughly two and a half hours per program per week — before transitions, before tantrums, before the twenty minutes the technician spends re-pairing after the second table-push, before the data collection that ten simultaneous programs demand from a technician who also has to actually teach. Spread across six programs, the same twenty-five hours puts every target above four hours weekly, enough for the massed, dense, fluency-oriented teaching that early repertoires are built from. Neither number is magic. But one of the two clinicians' did the division, and the plan shows it. When an experienced clinician looks at a goal list, the first thing she computes is not whether the goals are good. It is whether the goals divide into the hours.

Question two: what unlocks the most other learning? Both plans answer "manding and imitation" — the Goal Banks payroll logic, language earning its keep from day one — and on this the field genuinely agrees. But look one layer deeper at the manding goals and you find the chapter's finest-grained disagreement. Ramos' wrote "vocal approximation or sign," a standard disjunction that leaves the modality decision to the implementation team. Whitfield' wrote the decision itself into the goal: shape vocal inside the mand, guarantee the sign channel from day one, honor any consistent topography immediately. The difference looks typographical. It is not. Under Ramos's phrasing, a technician facing a frustrated Mateo at 4 p.m. on a Thursday makes the modality call.

Under Whitfield's, the analyst already made it, and made it visible enough to be argued with. A goal that leaves its hardest decision to whoever happens to be holding the clipboard has not finished being written.

Question three: what did the family actually ask for, and what did you do with it? This is where the chapter turns, because this is where the better plan on paper becomes the more dangerous plan in the world. The mother asked for two things: say Mama and look at me when I call him. Ramos's plan answers both, visibly — echoics on the list, name on the list. She can sit across the kitchen table and point to goals 2 and 4 and say, truthfully, "we are working on exactly what you asked for." Whitfield's plan postponed echoics for ninety days — for reasons that are, as it happens, clinically superior. Mand-first for this profile is the right call. The postponement list even says why. But here is what the intake documents show he did not do: he never translated the decision for her. The reasoning sat in the plan, in clinicians' language, in a document she signed but was never walked through. Three weeks into services, the mother asked the technician — not Whitfield', the technician — "when does he start working on talking?" and the technician, who had a plan but not the reasoning, said "I think thats later."

Run that sequence in slow motion, because every frame matters. The clinically stronger sequencing decision, communicated to no one, produced an alliance wound in week three; and an alliance wound in week three of a first authorization is not a soft-skills problem — it is a threat to every goal in the plan, because caregiver implementation is the multiplier under all of them and multipliers built on doubt fail quietly. What Whitfield' needed was one sentence, spoken at the kitchen table in the first week: "Everything in this plan is the road to Mama — requesting is how first words come, and if we drill sounds before he has a reason to make them, we slow the very thing you asked me for. In ninety days we check, and if the words arent starting to come through the requests, we go after them directly, and here is that date in writing." That sentence costs forty seconds. Its absence nearly cost the case. Postponement is a clinical decision; an uncommunicated postponement of the familys number-one priority is a clinical error, and it is an error committed disproportionately by strong clinicians', because they are the ones confident enough to postpone.

Question four: what did each clinician do with what they could not explain? Mateo responds to his name one time in five. Ramos' read this as a skill deficit and wrote a teaching program — reasonable; it usually is one. Whitfield' noticed something quieter in the packet: no hearing evaluation since the newborn screen, and a response pattern (best when the caller stands between him and a want) that is exactly what partial hearing or purely motivational control would each produce. He wrote the teaching program and the audiological referral, in parallel, and staged the goals criterion so that the first phase would generate data capable of telling the two stories apart. This is the move to study: he did not delay treatment to resolve the ambiguity, and he did not let treatment proceed as though the ambiguity did not exist. He built the differential into the program. Chapter 2 makes this a formal instrument; for now, notice only that one plan contains a hypothesis and the other contains a hypothesis test, and that from the outside they look nearly identical.

Question five — the one almost nobody writes down: what happens to everything you didnt choose? Ramos's plan has no postponement list, and so, in a sense that matters, Ramos's plan has no edges. The feeding log gestures at one deferral; the rest of the unchosen — haircuts, the car seat, listener expansion, sleeps behavioral half — simply isn't in the document, which means that at reauthorization no one, including Ramos', will be able to distinguish deliberately deferred from never considered. Whitfield's second page changes the entire epistemic status of his first page. Because the plan says what was postponed and why, every one of his

six inclusions is implicitly argued for — the goals that made the cut visibly beat named alternatives. Because each postponement carries a re-entry condition (day-90 echoic probe, post-medical sleep review, authorization-two car seat), the plan has a future built into it. And because the father's priorities appear on the second page with dates instead of disappearing, the family knows the difference between not yet and not heard. The postponement list is one page. It is the cheapest page in clinical writing, and it is the page that most reliably separates a plan that was reasoned from a plan that was assembled.

And one datum in the packet was weighed by neither clinician, recorded here on purpose. Mateo pushed the materials off the table — twice — and nothing in either plans pages says a word about what that push was. Both clinicians' read it as a behavior to program around; neither read it as a communication to answer. Hold the observation; it is not this chapters argument to make. Chapter 18 exists to give that push the hearing both plans didnt, and when you arrive there, remember that you saw it on the first day and moved past it too.

THE DEFINITION THIS BOOK RUNS ON

The trained reflex, arriving right about now, is to ask which plan is better — meaning which plan will produce the better outcome for Mateo. Refuse the question. It smuggles in an assumption this book will spend twenty chapters dismantling: that outcomes grade decisions. They do not. Mateos next twelve months will be co-authored by his neurology, his daycare, his grandmother's two afternoons, a funding decision no clinician controls, and a hundred variables that appear on no data sheet. A sound plan can meet a disengaging family; an assembled plan can meet an easy responder and look brilliant at reauthorization. Grade your decisions by their results and you will spend a career learning superstitions — repeating whatever happened to work, abandoning whatever happened to fail, and calling the residue experience.

So the standard has to attach to the decision at the moment it was made, with the information available when it was made — and you have already watched all three of its parts operate in this case. Name them now and they are yours.

A decision is visible when a stranger with your training, reading only the record, can reconstruct what you decided and why — including what you decided against. You saw the cost of its absence twice in one plan: Ramos's soundest judgments live nowhere but in Ramos', which means that for supervision, for reauthorization, for the BCBA who inherits the case, and for Ramos' herself in eight months, they barely exist.

A decision is evidence-linked when it cites the finding that licensed it — and the grain of the linkage is the tell. "Respond to name: B11" links a goal to a deficit. "Respond to name, staged criterion, audiology in parallel" links a goal to the ambiguity inside the deficit, which is why one of these programs will resolve a question this year and the other will only accumulate data about it.

And a decision is revisable when the conditions that would change it are stated in advance, with dates. Almost no treatment plan in this field has this property, and it is the one that most cleanly marks senior reasoning. Whitfield's echoic postponement is not a verdict; it is a scheduled experiment with its falsifier pre-registered — if correspondence emerges, or if the modality stalls, C1 enters the plan — which means that in ninety days nobody, including him, can quietly move the goalposts. A decision without stated revision conditions is not a strong decision. It is an unfalsifiable one, and unfalsifiable decisions are how goals end up "in progress" two authorizations later, un-met and un-questioned.

Visible, evidence-linked, revisable. That is the whole standard, and it is deliberately silent about outcomes. Run it against both plans and the chapters real lesson surfaces: Whitfield's document scores higher — and Whitfield' committed the cases most consequential error, because visibility on paper is not visibility to the family, and the standard applies to every audience a decision touches. The standard is not for crowning winners. It is an instrument for locating the specific weakness in any plan put in front of you — a skill you are about to need, because the next plan in front of you is yours.

NOW AUDIT YOUR OWN

Take out the plan you wrote before meeting Ramos' and Whitfield', and be clear about what you are checking. Not convergence — on the panel data you will meet later in this book, experienced clinicians' split on cases far cleaner than this one, and matching Whitfield' proves only that you and Whitfield' share defaults. You are auditing the plan against the standard, and the audit has teeth in four places.

First, the order. Did you write a sequence or a list? A list says these things matter. A sequence says I know what this authorization is for — and if a colleague reversed your goals three and five, would anything in your document object?

Second, the mother. "Say Mama" is either a program on your first page or a translated, dated deferral on your second. If it is on neither, you have just committed Whitfield's error with the advantage of never having met the family, and you now know its price. If you programmed echoics in month one, you have company — Ramos', most of the field, and possibly the correct answer for a Mateo whose echoic testing ran slightly warmer. The audit question is not whether echoics belong there. It is whether you chose them or defaulted to them, and the test of choosing is that you can name the data that would have made you postpone.

Third — the check this chapter has not asked yet — find the decision you delegated without meaning to. Somewhere in your plan is a phrase like "vocal or sign," "as tolerated," "when appropriate," "with supports as needed." Each one is a decision you exported to whoever holds the clipboard at 4 p.m. on the hard Thursday. Circle every instance. Some delegations are legitimate; a first plan usually contains two or three that are simply unfinished thinking wearing flexible language.

Fourth, the predictions. Now that you have watched two experts commit, notice whether you want to revise your percentages — and then notice what that urge is made of, because the two possible ingredients deserve different treatment. Expert judgment is evidence; if Whitfield's reasoning genuinely changed what you believe about Mateo, updating is calibration, not weakness. But confidence that moves merely because it saw company is the other thing, and you cannot tell which you have unless the record survives the update. So: leave the original numbers standing, and if you revise, write the new number beside the old with one line about what, specifically, moved you. Both columns come due in Chapter 20 — and the gap between them will be scored too.

BASELINE — TWELVE FORKS

Before this chapter closes, one more commitment, and this one runs the length of the book. Below are twelve clinical forks, each stripped to two lines. For each, commit to a branch — write the letter — and record your confidence as a percentage. No discussion follows. These are your baseline: the book will teach into every one of them across twenty chapters, and in Chapter 20 you will re-take the set, compare, and compute what changed. The forks are deliberately underspecified; when you find yourself wanting one more fact before choosing, write down which fact. Wanting the right missing fact is a skill this book measures too.

An FBA shows aggression at high rate in demand contexts — and equally high when demands are absent but a preferred adult is present. Lead hypothesis: (A) escape, multiply controlled / (B) attention, with demands as a discriminative context. Confidence: ____%

A break-mand program is working; problem behavior is down 80% at week six. The team wants to thin the break schedule. You: (A) thin now, ride the momentum / (B) hold two more weeks and define thinning criteria first. ____%

A fluent nine-year-old explains her coping plan flawlessly and has not used it once in three weeks of triggers. The primary problem is: (A) motivation / (B) the plan measured the wrong behavior from the start. ____%

Family requests toileting as goal one; child meets readiness markers; both parents work full-time and say weekends only. You: (A) launch a weekend-intensive protocol / (B) decline for now and program the prerequisites you can staff. ____%

Authorization returns at 60% of requested hours. You: (A) keep all goals, reduce dose per goal / (B) cut goals to preserve dose. ____%

An AAC learners device mands plateau; vocal approximations are appearing. The SLP wants to hold the device steady; a parent wants to drop it and "push speech." You: (A) hold both channels, honor everything / (B) run a structured comparison phase with the SLP and let the data pick the lead. ____%

A goal has been "in progress" for two authorization periods with a slowly rising trend. The likeliest problem: (A) insufficient dose or teaching error / (B) the goal itself — wrong target or missing prerequisite. ____%

A teenager tells you, mid-session and calmly, that a program you consider central "makes him feel like a dog being trained." You: (A) modify the programs form, keep its function / (B) suspend it and reopen the goal itself with him at the table. ____%

Your graphs show mastery; the technicians session notes keep mentioning prompts the data sheet has no column for. First move: (A) retrain data collection / (B) treat the mastery itself as unverified and re-probe cold. ____%

A caregiver has 40% protocol fidelity and rising burnout markers. The next caregiver goal targets: (A) fidelity, with more coaching / (B) the caregivers own behavior-sustaining repertoire — smaller protocol, deposits, respite — accepting slower client progress. ____%

A school placement decision hangs partly on your data, and your data are ambiguous. You: (A) present the ambiguity plainly and decline to tip the decision / (B) make a recommendation anyway, because declining is itself a decision with consequences. ____%

You inherit a colleagues case; the plan strikes you as significantly flawed, and the family loves it. First ninety days: (A) revise now, explain honestly / (B) run it mostly as written while your own assessment accrues standing to change it. ____%

Twelve answers, twelve percentages, dated. Some of these have answers the field considers settled; some split expert panels close to the middle, and part of what you will learn is which are which — knowledge the field has never systematically given its clinicians'.

FOR SUPERVISION

The swap. This chapter is built to be run, not just read. Supervisor and supervisee: write plans for Mateo independently, before either reads the chapters plans, then put the two documents side by side and diff them line by line. Every divergence is the curriculum — not who is right, but what question did we each answer differently. Run the five reconstruction questions from this chapter against both documents. The supervisor goes first in admitting an invisible decision; the exercise fails if the junior clinician is the only one being audited.

Three questions for the hour:

Find one decision in your current real caseload that exists only in your head — sound, defensible, and written nowhere. What would it cost to make it visible, and what has its invisibility already cost?

Whitfield's error was a translation failure on a correct decision. Which decision in your current plans is most in need of its forty-second kitchen-table sentence — and can you say the sentence, right now, out loud?

Pull one active plan from your caseload and hunt its delegated decisions — every "as tolerated," "when appropriate," "or sign." For each: who is actually making that call right now, at what time of day, under what conditions — and would you ratify the call they're making?

Companion references. Mateo's presentation is the Stack Books Profile 1 (Early Learner); Whitfield's plan is nearly the Early Learner Stack with one substitution — K7 in place of F13's joint-attention slot — made because this family's daily life was the iPad and the stacks own overlap logic permits exactly this trade. Ramos's plan draws A1, C1, C3, B11, B7, E1, K7, R1, P1, and the G8 watch-item from the Bank. The point of the comparison is not that the stack won; it is that the stack, individualized and argued, beat the same components assembled without argument — which is the only way a stack is ever supposed to win.

THE LETTER

Eleven days after both plans were submitted, the funding decision arrived. It authorized fifteen hours.

Not twenty-five. Fifteen — with a note inviting resubmission of intensity justification at ninety days, pending demonstrated progress.

Do the arithmetic before reading on. Ramos's ten goals now divide into ninety minutes each per week, which is not a teaching schedule; her plan, as written, no longer exists. Whitfield's dosage argument — the entire quantitative case for his six-goal design — was built on a number he no longer has; his plan, as written, no longer exists either. Every hour either clinician spent perfecting the document has just been converted into a different question, the question that was underneath the whole exercise all along: what, exactly, do you cut first, and how do you know?

One of these two clinicians' opened the letter and had to start reasoning from scratch. The other opened the letter, took out a page that already ranked every goal against every other goal with the reasons attached, and started crossing things off the bottom.

You have your own plan for Mateo in front of you, and now you have his real authorization. Before Chapter 2, revise your plan to fifteen hours — Chapter 3 will collect it, but do the cutting now, while the letter still stings. Notice which of your goals you fight to keep, and whether you can say why in a sentence a stranger could follow.

The plans were never the asset. They never are.

...